

We will try to simulate incident with one of the instances and see how AWS handles and how long it takes

As we can see around 20:18 both instances are working as expected and return the internal hostname, in this case

ip-10-0-1-46.eu-west-3.compute.internal

ip-10-0-2-94.eu-west-3.compute.internal

```
~ % date "+%H:%M:%S" && http http://app-lb-2017845830.eu-west-3.elb.amazonaws.com
20:18:20
HTTP/1.1 200 OK
Connection: keep-alive
Content-Length: 39
Content-Type: text/plain; charset=utf-8
Date: Thu, 03 Apr 2025 16:18:20 GMT
server: uvicorn
ip-10-0-1-46.eu-west-3.compute.internal

~ % date "+%H:%M:%S" && http http://app-lb-2017845830.eu-west-3.elb.amazonaws.com
20:18:20
HTTP/1.1 200 OK
Connection: keep-alive
Content-Length: 39
Content-Type: text/plain; charset=utf-8
Date: Thu, 03 Apr 2025 16:18:21 GMT
server: uvicorn
ip-10-0-2-94.eu-west-3.compute.internal
```

And the same goes for the health check. Both health checks are working, and both instances are healthy.

```
~ % date "+%H:%M:%S" && http http://app-lb-2017845830.eu-west-3.elb.amazonaws.com/healthcheck
20:20:22
HTTP/1.1 200 OK
Connection: keep-alive
Content-Length: 6
Content-Type: text/plain; charset=utf-8
Date: Thu, 03 Apr 2025 16:20:23 GMT
server: uvicorn

OK 200
```

Around 20:22, we introduced a fault in the system by triggering the /terminate-instance path, which changes the status code to 404 for the /healthcheck path. As you can see, after a few seconds, one of the hosts returns a "Healthcheck Fails" with 404 error code, while the other instance works.

```
~ % date "+%H:%M:%S" && http http://app-lb-2017845830.eu-west-3.elb.amazonaws.com/terminate-instance
20:22:10
HTTP/1.1 200 OK
Connection: keep-alive
Content-Length: 29
Content-Type: text/plain; charset=utf-8
Date: Thu, 03 Apr 2025 16:22:11 GMT
server: unicorn

Healthcheck status set to 404

~ % date "+%H:%M:%S" && http http://app-lb-2017845830.eu-west-3.elb.amazonaws.com/healthcheck
20:22:13
HTTP/1.1 404 Not Found
Connection: keep-alive
Content-Length: 21
Content-Type: text/plain; charset=utf-8
Date: Thu, 03 Apr 2025 16:22:13 GMT
server: unicorn

Healthcheck Fails 404

~ % date "+%H:%M:%S" && http http://app-lb-2017845830.eu-west-3.elb.amazonaws.com/healthcheck
20:22:15
HTTP/1.1 200 OK
Connection: keep-alive
Content-Length: 6
Content-Type: text/plain; charset=utf-8
Date: Thu, 03 Apr 2025 16:22:15 GMT
server: unicorn

OK 200
```

After a few moments, we can see that one of the instances in the target group became unhealthy and is draining. This means the Auto Scaling group detected the unhealthy instance and is creating a new one.

Target groups (1/1) [Info](#)

Filter target groups

< 1 > ⚙

✓

Name

▼

ARN

▼

Port

▼

Protocol

▼

Target type

▼

Load balancer


▼

VPC ID

▼

✓

[app-target-group](#)

 [arn:aws:elasticloadbalancin...](#)

5000

HTTP

Instance

app-lb

vpc-0f9242edb52ace8c4

Target group: app-target-group

Details

Targets

Monitoring


Health checks

Attributes

Tags

Registered targets (2) [Info](#)

Anomaly mitigation: **Not applicable**



Deregister

Register targets

Filter targets

< 1 > ⚙

☐

Instance ID

▼

Name

▼

Port

▼

Zone

▼

Health status

▼

Health status details

Administrative o...

▼

Override details


☐

[i-00e481f5f7ba61a49](#)


python-app-asg

5000

eu-west-3b (euw3-az2)

 Draining

Target deregistration is in progress

 No override

No override is current


☐

[i-08c5f7b4373c3f6a1](#)


python-app-asg

5000

eu-west-3a (euw3-az1)

 Healthy

-

 No override

No override is current

After a minute or so we can see a new instance created and healthy.

Target groups (1/1) [Info](#)

Filter target groups

app-target-group

arn:aws:elasticloadbalancin...

5000

HTTP

Instance

app-lb

vpc-0f9242edb52ace8c4

Target group: app-target-group

Registered targets (3) [Info](#)

Anomaly mitigation: Not applicable

Deregister

Register targets

Target groups route requests to individual registered targets using the protocol and port number specified. Health checks are performed on all registered targets according to the target group's health check settings. Anomaly detection is automatically applied to HTTP/HTTPS target groups with at least 3 healthy targets.

Filter targets

<input type="checkbox"/>	Instance ID	Name	Port	Zone	Health status	Health status details	Administrative o...	Override details
<input type="checkbox"/>	i-0c9d0c5f7b2469a9c	python-app-asg	5000	eu-west-3b (euw3-az2)	Healthy	-	No override	No override is curre
<input type="checkbox"/>	i-00e481f5f7ba61a49	python-app-asg	5000	eu-west-3b (euw3-az2)	Draining	Target deregistration is in progress	No override	No override is curre
<input type="checkbox"/>	i-08c5f7b4373c3f6a1	python-app-asg	5000	eu-west-3a (euw3-az1)	Healthy	-	No override	No override is curre

New host with internal host ip-10-0-2-173.eu-west-3.compute.internal

```
~ % date "+%H:%M:%S" && http http://app-lb-2017845830.eu-west-3.elb.amazonaws.com/
20:27:24
HTTP/1.1 200 OK
Connection: keep-alive
Content-Length: 40
Content-Type: text/plain; charset=utf-8
Date: Thu, 03 Apr 2025 16:27:24 GMT
server: uvicorn

ip-10-0-2-173.eu-west-3.compute.internal
```

Incident Summary

One instance in the autoscaling group failed its health check with a 404 response. The autoscaling mechanism terminated the unhealthy instance and launched a replacement, which passed the health check and functioned normally.

However, there are scenarios where the autoscaling group may not be able to heal the system:

- When S3 is unavailable
- When there is a configuration error in the Python file or in the AWS Terraform configuration
- When three availability zones are not available